

Solutions for Homework Assignment #1

Answer to Question 1. $T(n)$ is $\Theta(n)$. To prove this, we now show that $T(n)$ is both $O(n)$ and $\Omega(n)$:

1. $T(n)$ is $O(n)$.

This is not obvious because the procedure has two loops, one nested inside the other, each loop with an index ranging from 1 to n . So it may seem that the inner statement of line 4 is executed n^2 times (once for each value of i and j where $1 \leq i \leq n$ and $1 \leq j \leq n$). As we show below, however, this can never happen.

Claim: For every possible input array A , the **outer** loop of lines 2-5 is executed **at most 2 times**.

Proof: Let $n \geq 2$, and consider any input array $A[1..n]$.

If $A[1] \neq n - 1$ then `NOTHING(A)` returns in line 5 of iteration $i = 1$ of the loop of lines 2-5.

Now assume that $A[1] = n - 1$ (*).

If $A[2] \neq n - 2$ then `NOTHING(A)` returns in line 5 of iteration $i = 2$ of the loop of lines 2-5.

Now assume that $A[2] = n - 2$ (**).

From (*) and (**), $A[1] + A[2] = 2n - 3$.

So `NOTHING(A)` returns in line 5 of iteration $i = 2$ of the outer loop of lines 2-5.

Thus, in **all** possible cases, the **outer** loop of lines 2-5 is executed **at most 2 times**. **Q.E.D.**

The above claim implies the following:

For all $n \geq 2$, for every input A of size n , `NOTHING(A)` executes at most $2n$ iterations of the inner loop of lines 3-4.

Since each one of these inner loop iterations takes a **constant** time, it now clear that:

**There is a constant $c > 0$ such that for all $n \geq 2$,
for every input A of size n , `NOTHING(A)` takes at most $c \cdot n$ time.**

So $T(n)$ is $O(n)$.

2. $T(n)$ is $\Omega(n)$.

This is not obvious because the procedure may return “early” (e.g., after executing only a constant number of inner loop iterations) because of the loop exit conditions in line 4 and 5. Thus, to show that $T(n)$ is $\Omega(n)$, we must show that there is at least one input array A such that the procedure executes a linear number of inner loop iterations on this input, **despite the loop exit conditions of lines 4 and 5**. We do so below.

Let $n \geq 2$ and consider the input $A = [n, n - 1, n - 2, \dots, 2, 1]$, i.e., $A[n - j + 1] = j$ for all $1 \leq j \leq n$. Since $A[n - j + 1] \neq j$ does **not** hold for **any** j such that $1 \leq j \leq n$, with **this** input A , the **inner** loop of lines 3-4 **never** returns because of the condition $A[n - j + 1] \neq j$ in line 4.

So for **this** input A , `NOTHING(A)` executes **at least** n iterations of the **inner** loop of lines 3-4.

Since each one of these iterations takes a **constant** time, we conclude that:

**There is a constant $c > 0$ such that for all $n \geq 2$,
for some input A of size n , namely $A = [n, n - 1, \dots, 1]$, `NOTHING(A)` takes at least $c \cdot n$ time.**


Important note: For many arrays A of size n , for example all those where $A[n] \neq 1$, those where $A[n] = 1$ but $A[n - 1] \neq 2$, etc..., the execution of procedure `NOTHING(A)` takes only constant time! This is because the execution stops “early”, in line 4 (with $i = 1$), on these arrays.

So to prove that the worst-case time complexity of the procedure `NOTHING()` is $\Omega(n)$, a correct argument **must explicitly describe** some specific input array A of size n for which the execution of `NOTHING(A)` does take time proportional to n .

Answer to Question 2. The basic idea is to combine an n -bit vector with a complete binary tree, and use a structure that represents this tree without pointers.

a. The n -bit vector that represents S is as follows: the i -th bit of the vector V is set to 1 if and only if $i \in S$. The bits of V are also the leaves of a complete binary tree. Each node v of that tree is a single bit such that $v = 1$ if and only if the subtree rooted at v contains at least one leaf that is set to 1, i.e., it contains a leaf i such that $i \in S$.


For example, the data structure representation $n = 8$ and $S = \{1, 2, 6\}$ is as follows:



Note that each node v in the tree is just the logical “OR” of its two children.

To avoid pointers in the above tree structure (and thus save space), we use the *array representation* of a complete binary tree, just like we do to represent the complete binary tree of a HEAP! So we use a binary array A where bits $A[2i]$ and $A[2i + 1]$ are the left and right children of bit $A[i]$, respectively.

For example, the array A for $n = 8$ and $S = \{1, 2, 6\}$ is as follows:



This data structure takes only $2n - 1$ bits of space: n bits for the n -bit vector V representing S , and $n - 1$ bits for the nodes of the tree above V .

Note that the entire structure can be thought of a *single* complete binary tree with $2n - 1$ nodes: n leaves and $n - 1$ internal nodes. The depth of this tree is $\log_2 n$. In our example, $n = 8$ and the depth is $\log_2 8 = 3$.

b. MAXIMUM: the maximum element of S is the index of the rightmost “1” in the bit vector V (which is the also the rightmost leaf that contains 1 in the complete binary tree). To find the rightmost “1” in V , start at the root of the complete binary tree, and follow the rightmost path that contains only 1’s. In the example above, this path is $A[1], A[3], A[6], A[13]$, and it leads to the 6-th element in the binary vector, i.e., it leads to element 6 of S .

The *index* of the rightmost “1” in the n -bit vector is just its index in the array A minus $n - 1$. In the example above, the rightmost “1” in the 8-bit vector is $A[13]$, and its index in the n -bit vector is $13 - (n - 1) = 13 - 7 = 6$. Note that 6 is indeed the maximum element of the set S in our example.

The worst-case time complexity of the MAXIMUM operation is proportional to the depth of the tree, more precisely it is $\Theta(\log_2 n)$.

INSERT(j): Set bit j of the n -bit vector to 1, and set all its ancestors in the tree also to 1. More precisely, start from the bit A that represents element j of S , namely bit $A[j + (n - 1)]$, and set it to 1. Then go to its parent $A[\lfloor j + (n - 1)/2 \rfloor]$ and set it to 1, etc., and continue this way up to the root and set $A[1]$ to 1.

Alternatively, one can start by setting $A[1]$ to 1, and do the same for every element of A in the path from $A[1]$ to $A[j + (n - 1)]$ (think about how you would find this path...).

The worst-case time complexity of an INSERT(j) operation is $\Theta(\log_2 n)$.

c. MEMBER(j): Check whether bit j of the n -bit vector is 1. More precisely, MEMBER(j) = TRUE iff $A[j + (n - 1)] = 1$. The worst-case time complexity of this operation is clearly $\Theta(1)$.