

Solutions for Homework Assignment #2

Answer to Question 1.

Whiz Warehouses needs can *mostly* be satisfied by using max binomial heaps to implement mergeable priority queues. Each warehouse stores its set of orders in a separate max binomial heap, where for order x , $\text{key}[x]$ is the integer amount of x 's incentive. You can find descriptions of the operations on ordinary binomial heaps in the CLRS (2nd Edition) hand-out in course materials. The one operation that this approach will not satisfy with the required performance guarantee is finding the current maximum incentive in constant time: MAXIMUM on a max binomial heap with n nodes has worst-case $O(\log n)$ running time.

Provide constant-time maximum incentive look-up by augmenting a max binomial heap H with a field max to record the current maximum incentive. The augmented data structure is simply $H_A = (H, \text{max})$.

In what follows assume that $\text{key}[\text{NIL}]$ returns special value $-\infty$. The operations required by Whiz Warehouses are below, with a subscript "A". Each operation's correctness follows from the correctness of the component binomial max heap operations, and from the properties of the **max** function, in particular binomial heap operation MAXIMUM is called to find a maximum priority after operations that may rearrange the heap.

- **Make-Heap_A**()
 - 1 allocate a new object H_A
 - 2 $H_A.H \leftarrow \text{MAKE-HEAP}()$
 - 3 $H_A.\text{max} \leftarrow -\infty$
 - 4 **return** H_A

WORST-CASE TIME COMPLEXITY:

The new operation inherits worst-case time complexity $O(1)$ from MAKE-HEAP.

- **Extract-Max_A**(H_A)
 - 1 $\text{tmp} \leftarrow \text{EXTRACT-MAX}(H_A.H)$
 - 2 $H_A.\text{max} \leftarrow \text{key}[\text{MAXIMUM}(H_A.H)]$
 - 3 **return** tmp

RUNNING TIME:

Calling one procedure with $O(\log n)$ worst-case time complexity and then another with $O(\log n)$ worst-case time complexity produces one with $O(\log n)$ worst-case time complexity.

- **Union_A**(H'_A, H''_A)
 - 1 $H_A \leftarrow \text{MAKE-HEAP}_A()$
 - 2 $H_A.\text{max} \leftarrow \text{maximum of } H'_A.\text{max} \text{ and } H''_A.\text{max}$
 - 3 $H_A.H \leftarrow \text{UNION}(H'_A.H, H''_A.H)$
 - 4 **return** H_A

RUNNING TIME:

The new operation inherits worst-case time complexity $O(\log n)$ from UNION, where n is the sum of the number of nodes in H' and H'' , plus worst-case constant time complexity to calculate and update $H_A.\text{max}$. Thus the worst-case time complexity is $O(\log n)$.

- **Insert_A**(H_A, x)
 - 1 $\text{INSERT}(H_A.H, x)$
 - 2 $H_A.\text{max} \leftarrow \text{maximum of } H_A.\text{max} \text{ and } \text{key}[x]$.

RUNNING TIME:

The new operation inherits worst-case time complexity $O(\log n)$ from Insert, plus worst-case constant time to access and update $H_A.\text{max}$, so the result has worst-case time complexity $O(\log n)$.

- $\text{Max-Incentive}_A(H_A)$
1 **return** $H_A.\text{max}$

RUNNING TIME:

The new operation has worst-case time complexity $O(1)$.

Answer to Question 2.

The binary representation of 56 is $\langle 111000 \rangle_2$. Thus the 56-node binomial heap has three binomial trees: one with 2^5 nodes, one with 2^4 nodes, and the smallest with 2^3 nodes (note that these correspond to the 1s in the binary representation of 56). We know that each binomial tree with 2^k nodes, has height k , and the degree of its root is k , so the degree of the smallest tree's root is 3. After an EXTRACT-MIN operation the heap holds 55 nodes, and the binary representation of 55 has a "1" in the right-most position. So the smallest tree has $2^0 = 1$ node and its degree and height is 0.

Answer to Question 3.

ALGORITHM:

Search for the maximum key in B_1 and the minimum key in B_2 *in parallel* and *stop as soon as one of them is found*, as follows. Traverse down the rightmost path of B_1 and down the leftmost path of B_2 , by following the right-child pointers in B_1 and the left-child pointers in B_2 , *in parallel*, i.e., by alternating steps; stop this traversal as soon as you find a node x such that:

(*) x is in B_1 and x has no right child, **or** x is in B_2 and x has no left child.

Remove x from its tree (note that this takes only constant time as x is simply replaced by its only child), and then make x the root of the new, merged tree T with b_1 (the root of B_1) as its left child and b_2 (the root of B_2) as its right child.

CORRECTNESS:

From the ordering of keys in a binary search tree, the key at x is either the largest key of B_1 or the smallest key of B_2 . Since every key in B_1 is smaller than every key in B_2 , every key in B_1 is smaller than or equal to the key of x , and every key in B_2 is greater than or equal to the key of x .

Since: (a) each of B_1 and B_2 are binary search trees, and (b) the keys in x 's left subtree are smaller than or equal to the key of x , and the keys in x 's right subtree are greater than or equal to the key of x , the new tree rooted at x is also a binary search tree.

RUNNING TIME:

Note that the parallel traversals from the roots of B_1 and B_2 to the *first* node x that satisfies condition (*) can take at most $2 \cdot \min\{h_1, h_2\}$ steps. Removing x takes $O(1)$ time because it has only one child. Creating the new tree T and assigning the root's children also takes $O(1)$ time. Therefore, in total, the algorithm runs in $O(\min\{h_1, h_2\})$ time in the worst-case.