

Solutions for Homework Assignment #3

Answer to Question 1.




Figure 1: After all the insertions




Figure 2: After deleting 19

Answer to Question 2.

a. The data structure D is represented as an augmented AVL tree. Each node u of the augmented AVL tree contains information about a house listing. It contains fields $price(u)$, giving the price of the listing; $area(u)$, giving the floor area of the listing; and $maxarea(u)$ containing the maximum floor area of all the listings in the subtree rooted at u (including u). Field $price$ is used as the key of the node — thus, an in-order traversal of the tree visits all the listings in non-decreasing price order. In addition, each node u contains the usual information of an AVL tree node: pointers to the left and right children as well as the parent, and the balance factor.

b. The operations are implemented as follows:

- $INSERT(D, x)$: If x is a pointer to a new listing, insert the listing pointed to by x into D .

- (1) Use the ordinary BST insertion algorithm to insert x into the tree using the $price$ field as the key. (Note that it does not matter if the key of x already exists in the BST: if, while entering x , we encounter a record y with the same key as x , by we can consider x to be an immediate successor of y , and so we can continue the insertion of x in the right subtree of y .) Node x is now a leaf.
- (2) Traverse the path from the new leaf x to the root of the tree. For each node u along that path, update the $maxarea$ field according to the following identity:

$$maxarea(u) = \max(area(u), maxarea(lchild(u)), maxarea(rchild(u)))$$

(where we assume that $maxarea(NIL) = -1$).

- (3) Traverse the path from the new leaf x to the root of the tree, performing rotations and updating the balance factors as required by the ordinary AVL insertion algorithm. In addition, we update the $maxarea$ field of every node u involved in a rotation, according to the above identity.

We stop rotations and updates to the balance factors as in the ordinary AVL insertion algorithm.

Remark: The two traversals described above can be done in single pass where we first update the $maxarea$ field and we then do the balancing part (which may also require a few $maxarea$ field updates).

The time required by this algorithm is that required by the ordinary AVL insertion algorithm, i.e., $O(\log n)$, plus the time required to update the *maxarea* fields of the new node's ancestors, and the *maxarea* fields of the nodes involved in a rotation (a few nodes for each rotation). Since there are $O(\log n)$ such nodes and each *maxarea* update requires $O(1)$ time, the additional time is also $O(\log n)$. Thus, $\text{INSERT}(D, x)$ takes $O(\log n)$ time.

- $\text{DELETE}(D, x)$: If x is a pointer to a listing in D , remove that listing from D .
 - (1) Use the ordinary BST deletion algorithm to delete x from the tree using the *price* field as the key. Since the tree is an AVL tree, the deletion of x ultimately results in the removal of a leaf from the tree. Let z be the parent of that leaf.
 - (2) Traverse the path from z to the root of the tree. For each node u along that path, update the *maxarea* field according to the following identity:

$$\text{maxarea}(u) = \max(\text{area}(u), \text{maxarea}(\text{lchild}(u)), \text{maxarea}(\text{rchild}(u)))$$

(where we assume that $\text{maxarea}(\text{NIL}) = -1$).

- (3) Traverse the path from z to the root of the tree, performing rotations and updating the balance factors as required by the ordinary AVL insertion algorithm. In addition, we update the *maxarea* field of every node u involved in a rotation, according to the above identity.

We stop rotations and updates to the balance factors as in the ordinary AVL insertion algorithm.

Remark: The two traversals described above can be done in single pass where we first update the *maxarea* field and we then do the balancing part (which may also require a few *maxarea* field updates).

The time required by this algorithm is that required by the ordinary AVL insertion algorithm, i.e., $O(\log n)$, plus the time required to update the *maxarea* fields of the removed leaf's ancestors, and the *maxarea* fields of the nodes involved in a rotation (a few nodes for each rotation). Since there are $O(\log n)$ such nodes and each *maxarea* update requires $O(1)$ time, the additional time is also $O(\log n)$. Thus, $\text{DELETE}(D, x)$ takes $O(\log n)$ time.

- $\text{MAXAREA}(r, p)$: If r is the root of a subtree of D , return the largest floor area, amongst all listings in this subtree whose price is $\leq p$; if there is no such listing, then return -1 .

```

MAXAREA( $r, p$ )
  if  $r = \text{NIL}$  then
    return  $-1$ 
  else if  $p < \text{price}(r)$  then
    return MAXAREA( $\text{lchild}(r), p$ )
  else
    return  $\max(\text{maxarea}(\text{lchild}(r)), \text{area}(r), \text{MAXAREA}(\text{rchild}(r), p))$ 
  end if

```

Note that $\text{MAXAREA}(D, p)$ is simply $\text{MAXAREA}(r, p)$ where r is the root of tree D .

The time complexity of $\text{MAXAREA}(r, p)$ is proportional to the height of the AVL subtree of D rooted at r . (This is because (i) each call at a node u results in a *single* recursive call, at the left or the right subtree of u ; and (ii) each call involves a constant amount of work other than the recursive call that it makes.) If D contains n listings, the height of any subtree of D is $O(\log n)$, so the time complexity of $\text{MAXAREA}(r, p)$ is also $O(\log n)$.

Answer to Question 3.

We want to determine whether there is an $x_i \in X$ and a $y_j \in Y$ such that $x_i^2 + y_j^2 = d^2$, i.e., whether $\exists i, 1 \leq i \leq n$, and $\exists j, 1 \leq j \leq n$, such that $x_i^2 = d^2 - y_j^2$. Note that since every $x_i \in X$ and $y_j \in Y$ is an integer, $x_i^2 = d^2 - y_j^2$ can hold only if d^2 is also an integer.

a. Algorithm:

To do so we use a hash table T of size m and with a hashing function h (more about m and h later).

For each $i, 1 \leq i \leq n$:

INSERT x_i^2 in table T (in front of the list starting at $T[h(x_i^2)]$) /* each INSERT takes $O(1)$ time */

For each $j, 1 \leq j \leq n$:

SEARCH for $(d^2 - y_j^2)$ in table T (in the list starting at $T[h(d^2 - y_j^2)]$) /* this takes $O(\alpha)$ expected time */

If found then output “YES” and stop /* found an x_i such that $x_i^2 = d^2 - y_j^2$ */
else output “NO”

b. Assumptions:

1. SUHA (Simple Hashing Uniform Assumption): For each $i, 1 \leq i \leq n$, integer x_i^2 is equally likely to hash into any of the m slots of T , independently from where $x_1^2, x_2^2, \dots, x_{i-1}^2, x_{i+1}^2, \dots, x_n^2$ hash into.
2. The size m of T is “proportional” to n , more precisely m is $\Theta(n)$ (actually, $m \in \Omega(n)$ suffices).
3. Computing the hash function h takes $\Theta(1)$ time.

c. The expected running time is $\Theta(n)$:

- Each “INSERT x_i^2 in table T ” takes $\Theta(1)$ time. So inserting $\{x_1^2, x_2^2, \dots, x_n^2\}$ into T takes $\Theta(n)$ time.
- By Assumption 1, the expected length of each chain of T is $\alpha = n/m$. By Assumption 2, $\alpha = \Theta(1)$.
- So the expected running time for *each* “SEARCH for $(d^2 - y_j^2)$ in table T ” is $\Theta(1)$, and the expected time to do this search for *all* $j, 1 \leq j \leq n$, is $\Theta(n)$.

d. The worst-case running time is $\Theta(n^2)$:

1. It is $O(n^2)$ because:

- As explained above, the first loop of the algorithm (which inserts $x_1^2, x_2^2, \dots, x_n^2$) takes $O(n)$ time.
- Now consider the second loop of the algorithm. Since no chain in T can contain more than n elements, each “SEARCH for $(d^2 - y_j^2)$ in table T ” in this loop takes at most $O(n)$ time. Since the loop does at most n such searches, the loop takes at most $O(n^2)$ time.

2. It is $\Omega(n^2)$ because:

The following “time-consuming” execution can occur.

(i) All the integers $x_1^2, x_2^2, \dots, x_n^2$ hash into the same slot $T[k]$ of T (i.e., they all “collide” into $T[k]$). So they form a chain of length n starting in slot $T[k]$.

(ii) For every $j, 1 \leq j \leq n$, $(d^2 - y_j^2)$ also hashes into slot $T[k]$, but there is no i such that $x_i^2 = d^2 - y_j^2$. So every “SEARCH for $(d^2 - y_j^2)$ in T ” in the second loop of the algorithm “fails”: each search traverses the whole chain of length n at $T[k]$ and takes $\Omega(n)$ time. Thus the n searches for $(d^2 - y_j^2)$ in T take $\Omega(n^2)$ time.

A similar algorithm that also works:

For each $i, 1 \leq i \leq n$:

INSERT $d^2 - x_i^2$ in table T (in front of the list starting at $T[h(d^2 - x_i^2)]$)

For each $j, 1 \leq j \leq n$:

SEARCH for y_j^2 in table T (in the list starting at $T[h(y_j^2)]$) /* found a y_j such that $y_j^2 = d^2 - x_i^2$ */

if y_j^2 is found then output “YES” and stop
else output “NO”

Another algorithm that works:

For each $i, 1 \leq i \leq n$:

INSERT x_i in table T (in front of the list starting at $T[h(x_i)]$)

For each $j, 1 \leq j \leq n$:

$x := \sqrt{d^2 - y_j^2}$

SEARCH for x in table T (in the list starting at $T[h(x)]$)

if x is found then output “YES” and stop /* found an x_i such that $x_i = \sqrt{d^2 - y_j^2}$ */
else output “NO”
